UNIVERSITY OF WATERLOO

Mechanical and Mechatronics Engineering

DEVELOPMENT OF AN IOS APPLICATION

Crouton Labs Incorperated

Kitchener, ON

Prepared by

Jake Nielsen

ID #20338042

jake.k.nielsen@gmail.com

2B Mechatronics Engineering

May 15, 2012

Jake Nielsen
40 Karch st.,
Cambridge, ON N3C 1Y5

May 15, 2012

Dr. Pearl Sullivan, Chair
Mechanical and Mechatronics Engineering Department,
University of Waterloo,
Waterloo, ON N2L 3G1

Dear Dr. Pearl Sullivan:

Re: Submission of my work term report.

I have just completed my fourth work term, following my 2B term. Please find enclosed my first work term report entitled: "Development of an iOS Application" at Crouton Labs Incorperated.

I am one of the co-founders of Crouton Labs Inc. We make smartphone apps for restaurants in the Kitchener/Waterloo area. In particular, this report discusses the development choices that I made as the iOS developer on our development team while designing and programming the app that we made for Pita Factory.

I have had no direct assistance from anyone. I do wish to thank Christophe Biocca for his invaluable support in learning and working with the objective-c language. I can say without a doubt that my knowledge of the language, its features, and coding practices would not be half of what they are today if not for Christophe's expertise.

This report was written entirely by me and has not received any previous academic credit at this or any other institution.

Yours sincerely,

Jake Nielsen, 20338042

Encl.

Table of Contents

List of Figures

Summary

The purpose of this report is to explain the highlights of the software development project that I was involved with at Crouton Labs Inc. The project itself was meant to solve the problem of long line-ups at restaurants, and efficiency of fast food procuction lines. The report explains the coding practices that were used during the development process, the challenges that were faced during development, and the solutions that were implimented to solve those problems. This report emphasises the importance of using proper design patterns and coding practices by demonstrating the amount of time saved and the quality that can be achieved by adhering to these practices. The final product is one that solves the problem it was intended to solve, is stable, looks good, is optimized enough that it feels smooth to interact with, and has very readable code.

1 Introduction

The purpose of this report is to convey the challenges that were faced during the development of the iOS application that Crouton Labs Inc. provides to Pita Factory.

In order to undestand the motivation for writing this app in the first place, it's important to understand what problem this app is intended to solve. This app is an attempt to reduce the time spent in fast food restaurants waiting in line and processing transactions. This should allow for the restaurant to service more customers and thus increase profits. With this in mind, this main goals of this app are to allow a user to browse the restaurant's menu, construct an order that they would like to purchase, and to provide a means of processing the transaction through the app.

The programming language that was used to accomplish this goal is called objective-c. Objective-c is syntactically similar to c and smalltalk. Its strange syntactic combination can be confusing at first, as there are often two syntactically different ways of achieving exactly the same effect. Objective-c at first glance does not appear to be much different from c++ in terms of language features, but upon further inspection it can be seen that objective-c is perhaps better compared to a dynamic language. While objective-c does not fit the true definition of a dynamic language, it has many language features that are often associated with dynamic languages. Object and class methods can be called by progamatic string construction of the name of the method, which is refered to as a "selector". Similarly, subclasses can be refered to via progamatic strings. With some work, classes can even request the names of their subclasses, allowing for some very powerful programming techniques that will be discussed later in this report.

This report will discuss the general coding practices that were adhered to during development, the challenges faced during development, and the solutions that were implimented to deal with the challenges that arose.

The actual interface flow is not within the scope of this report. The user interface was designed by a 3rd party designer and was merely implimented and occaisionally extended upon by Crouton Labs.

2 Coding Practices

2.1 Motivation

The coding practices that were maintained during the development process are an extremely important factor that affected the final product of this development process. Proper coding practices are often dismissed or overlooked by many programmers to the detriment of their productivity and positive outlook. Adhering to good practices make the painful parts of software development much less frequent and much more bearable.

2.2 Version Control

Using version control is the single, most important practice any good programmer should have. While the choice of version control software is hotly debated, I am partial to git. Git is a very powerful source control system and is also very widely used. Part of good source control use is having a remote repository. For this project, www.github.com was used as our remote repository.

All version control software has the concept of small, incrimental changes that, when put together, move the development of the code-base forward. In git, these are refered to as commits. It's good practice to make very small, bite-sized commits. In escence, each commit should fix one bug, impliment one feature, or add one more logical piece to the codebase. In this way, if certain features aren't working, or simply aren't finished, they can be excluded from release builds without any painful merge conflicts.

If a feature is not trivial to impliment, the feature should be made on a separate "branch" soas to maintain the small size of commits while still making sure that it is known that the group of commits are associated (and likely dependant on) one another. In this way, if a feature needs to be included or discluded, this can be done simply by including, or discluding the entire branch.

2.3 Efficientcy vs Simplicity

This is another widely overlooked and undervalued concept that helps immensely during the development process. Many programmers sacrifice simplicity in favor of efficiency much more often than they should. While it's important under certain circumstances to have efficient code, it can be hard to know where optimization would do the most good, or even if it is necessary. Optimization should only happen if it becomes a problem. Overly optimized code is indistinguishable from properly optimized code at runtime, but at development time, overly optimized code is much less readable, and will cause pain for any contributors, or even for the person who did the optimization in the first place. Hence, optimize only when necessary.

2.4 Refactoring

It seems obvious, but refactoring is esscencial to maintaining a usable and agile codebase. To emphasize the importance that needs to be put into refactoring, the codebase for this project has been rewritten little-by-little about 3.3 times. That is, for every line of code in the codebase, there were 3.3 lines added and deleted. The general mantra that describes good refactoring is, "refactor early, refactor often".

2.5 Design Patterns

Design patterns are widely used methods of solving particular problems that can come up very frequently in programming problems. It's good to have a large arsenal of design patterns to call on when developing software. By using patterns that are widely known, your code will be more readable to someone who has never seen your code before, but has experience with design patterns. Knowlege of design patterns can drastically increase the speed at which you arrive at solutions. This is an area that was not adhered to as rigidly as the other proper coding practices, but it is something that would have added a lot of value. The aggressive refactoring resulted in a codebase that makes proper use of design patterns, but it would have been better to have arrived at them initially.

3 Challenges and Solutions

3.1 Section Structure

In this section, the challenges that were faced will be introduced and explained. The solutions that were implimented to solve these problems will also be discussed. It's worth noting that when the software was being developed, the approach that was taken involved itemizing all of the challenges that were expected to come up so that all of them could be taken into consideration before writing a single line of code. In general, it's good to catch the challenges and make sure that strategies for dealing with them can be incorporated nicely into the code base. The earlier a challenge is spotted, the faster and more effectively it can be solved.

3.2 Code Testing

In order to be certain that a software solution is relatively bug-free and stable, it's often good practice to write unit tests and system tests to make sure that the software operates as expected. In the case of an iOS application, this becomes a very difficult thing to do. Most of the bugs that occur in iOS applications are visual or system-behavioural. Since there is no straightforward way to write system tests for user interfaces, none were written in this development process. Extensive manual testing was done prior to releases to make sure that no breaking bugs were introduced to the live codebase. The lack of unit tests or automated system tests meant that all of the other coding practices discussed in the previous section became extremely important. The manual testing also needed to be extremely thorough to make absolutely sure that no bugs were pushed to the general public. While somewhat frowned upon, this method proved to be successful, however without almost obsessive adherence to proper coding practices it likely would not have worked very well. As it stands, only one release had a bug that was considered to be breaking, and was a relatively subtle behavioural bug. In terms of stability, to date, there has never been a crash in a production build.

3.3 Object Serialization

In order to serialize an object to save it to non-volitile memory, serialization and deserialization methods need to be implimented for each class that needs to be saved. There are 18 different classes within the codebase that need to be serializable. This quickly became tedious. In order to solve this problem, a class was created that utilizes some of the more powerful language features of objective-c. In particular, objective-c allows classes to programatically reference the values and the names of their member variables. By utilizing that feature, the class (called AutomagicalCoder) is able to cycle through all of its own member variables and store serialize them by encoding them with their own string name as the key. In this way, any class that subclasses AutomagicalCoder needn't worry about how to serialize itself, and instead just inherently knows how.

3.4 User Interface Tools

There are a lot of similarity between facets of a user interface. Particularly in this application, there was often need to have a way of displaying a multitude of different types of objects as cells in a table. Similarly, the behaviour invoked when the user taps on these object cells with his/her finger is often (although not always) the same. Both of these problems lent themselves well to the dispatcher design pattern.

In the dispatcher pattern, when it is necessary to do a similar action to many different types of data, instead of coding each action separately, the data is sent to a "dispatcher" that looks through the data that it knows how to handle, and decides accordingly which special actions to take depending on the data. If the data is not of a known type, depending on the dispatcher, it is either correct to take some default action, or simply to throw an error to let the programmer know that an unexpected piece of data was passed to the dispatcher. Objective-c lends itself to the dispatcher method, because of the language feature that allows a class to look through it's own subclasses without the subclasses explicitly making themselves known to the parent class.

Without delving any deeper, a solid working knowledge of Apple's API is required to understand most, if not all of the following paragraphs. An abbreviated explaination of Apple's API can be found in Appendix A.

In the case of displaying cells, the parent class (CustomViewCell) looks through it's subclass and calls the canDisplayData method to query them about the data that it has been presented with. This is shown in the figure below.

Figure 1: Illustration of the Dispatcher Design Pattern

In this way, all that is necessary to fully define and display a list view is a list of the data that should be displayed in the list. By dispatching to CustomViewCell, the particularities of displaying each individual cell needn't be handled in any way by the tableViewDelegate.

Another language feature involved in the simplification of the TableView interface provided by Apple is that of calling methods by constructing their name through string manipulation. If there's somthing that needs to occur when the user clicks on a cell, the TableViewDelegate need only impliment a function named ¡CELLNAME¿Handler. The method (if it exists) will be called by the table view delegate whenever the cell is clicked on. In this way, the delegate can distinguish between cell clicks without needing to actually know the type of the cell or the position of the cell (which is the traditional way that Apple intends for cell distinctions to be made).

By making all of these simplifying abstractions, the code length of the files defining the behaviour of particular pages is vastly. One class in particular (OrderComboViewController) has 4 lines of code in its header file and 23 lines of code in its implimentation. A screenshot of this particular page is shown below.

The UITableViewDataSource for this page is simplified in a similar way. The class responsible for the page contents (OrderComboRenderer) has 7 lines in its header file, and 37 in its implimentation.

The dispatcher pattern is a very powerful way of simplifying code.

Figure 2: Screenshot of the Order Combo Page
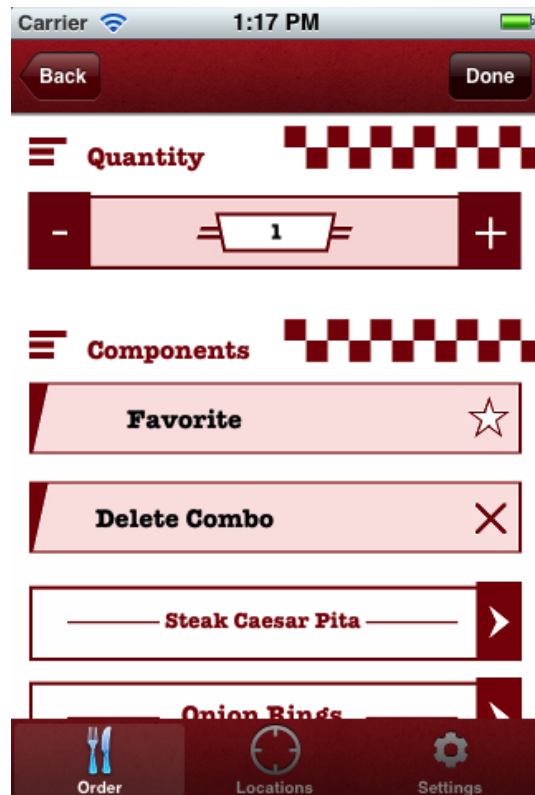
3.5 Client Version Management

One of the complicating factors in designing any application is how various versions of the code will interact with one another. The two possible modes of interaction in this case are communication with the server, and loading data cached by an older version of the client code.

3.5.1 Server Communication

If an old version of the client code were to try to interact with the server, the server's behavior may have to be different than if an up to date client app were to attempt communications. The way that this is solved is again by dispatching. As part of every communication with the server, the iOS application transmits its version number. Based on the version number, the server dispatches the the response behaviour to the appropriate version of the server code.

3.5.2 Client Data Caching

A more difficult complication is that of loading cached data. If the ipod/iphone user has an old version of the app, saves some data, updates their app to the current version, and then attempts to load the cached data, bad things can happen. Depending on the version of the app that was used to cache the data, the procedure to update the data to be consistant with the current model is different. To solve this problem, a combination of the dispatcher pattern and the fallthrough pattern was used. The first step in determining what to do with the data being loaded is to check the version string. Luckillty, this problem was anticipated, and the cached data always contains a version string. Because this all needs to be done by AutomagicalCoder, no actual knowledge of what these variables are is known explicitly at load time. To solve this, Automagical Coder checks the version string, and if the version string is not the current version string, it checks for the existance of a recovery routine for each of the variables that it determines that it needs to load. if a recovery routine is implimented in the child class, AutomagicalCoder dispatches to the recovery routine by calling ¡VARIABLENAME¿Recovery. The recovery routine uses the fallthrough pattern.

The fallthrough pattern uses a switch case structure where the case clauses do not contain a break statement. In this way, the variable information will be updated in a cascade through all of the versions starting at the version that was found on the disc all the way up to the present model. In this way, the only thing that needs to be updated from version to version is the logic for moving the data from the version immediately prior to the current version. Everything else just cascades through.

4 Conclusion

In conclusion, by using the habits and patterns discussed in this report, the development project was a great success. From a technological view point, this product turned out great. It's stable, it looks good, it is optimized enough that it feels smooth to interact with, and its code is very readable.

The app can currently be downloaded from the app store, and is called "Pita Factory". As not everyone has an iOS device, below are some screenshots of the finished product.

Figure 3: Screenshot of the Order/Menu Page

Figure 4: Screenshot of the Credit Card Page

Figure 5: Screenshot of the Location Page

A   Apple API Abbreviated

Appendix A is an abbreviated description of Apples's iOS API. It will explain the general concepts behind displaying things on an iOS device as well as the specific elements of the API that are pertainant to this report.

As a general overview, there are views, models, and controllers. The views represent the actual images that the user sees. The models represent the data that is meant to be conveyed to the user. The controllers encapsulate the ways that the user can interact with the data. This is known as the model-view-controller design pattern and is the design pattern that apple intends for programmers to use when interacting with their API.

A.1   UIViews

UIView is the most basic class that is meant to be displayed to the screen. A UIView has an area of the screen that it can manipulate via its drawrect method. A UIView can also have an arbitrary number of subviews which draw themselves onto the parent UIView. There are many built-in subclasses of UIView that can accomodate most applications that a person might want to build.

Some notable subclasses of UIView are: UILabel, UIButton, UITableView UIViewCell, and UIImageView.

A.1.1   UITableView

A UITableView is a view that can display any number of scrollable UIViewCells. The UITableView must be associated with a UITableViewDelegate in order to define any interactive behaviour and must be associated with a UITableViewDataSource in order to have any content. they will be discussed more in the UIViewControllers and UIDataSources sections.

A.1.2   UIViewCell

A UIViewCell is a UIView that can be produced by a UITableViewDataSource. The vanila UIViewCell can be used by itself as elements to populate a UITableView and can be configured in several different default configurations

that can accomodate most data that a person might want to display reasonably well. In cases where the vanila UIViewCell will not suffice, UIViewCell can be subclassed to accomodate arbitrarily complex and unusual cells.

A.2   UIViewControllers

UIViewController is a class that is not intended to be used by itself, but rather to be subclassed to accomodate specific behaviour that a programmer might intend for a view to have. View controllers are often made to be the delegates of the view that they control so that a view can request behavior from its view controller without actually needing to know anything specific about the view controller.

A.2.1   UITableViewDelegate

The most notable type of delegation that this report refers to is that of the UITableViewDelegate. The UITableViewDelegate is responsible for telling the UITableView what needs to happen upon the selection of a cell at a particular position on the table. It's also responsible for letting the UITableView know the height that each cell is intended to occupy.

A.3   UIDataSources

In general, the "model" part of the model-view-controller pattern is satisfied by a class that impliments a datasource interface defined for the UIView that it is acting as the model for. This is not always the case, and in cases where it isn't, the model will simply be a subclass of NSObject, which is a base class that everything derrives from within the the Apple API.

A.3.1   UITableViewDataSource

A class that impliments the UITableViewDataSource interface will provide a UITableView with the cells that the UITableView will populate itself with as well as letting it know the number of cells that exist in the table, the titles of the sections that make up the table.